

Deep Neural Network for MNIST Classification

Name : Ha Thanh Nga

Class : BA62

Subject : Big Data Analytics



Contents

01

Data Import &
Preprocessing

02


Batching &
Shuffling

03

Training the
Model

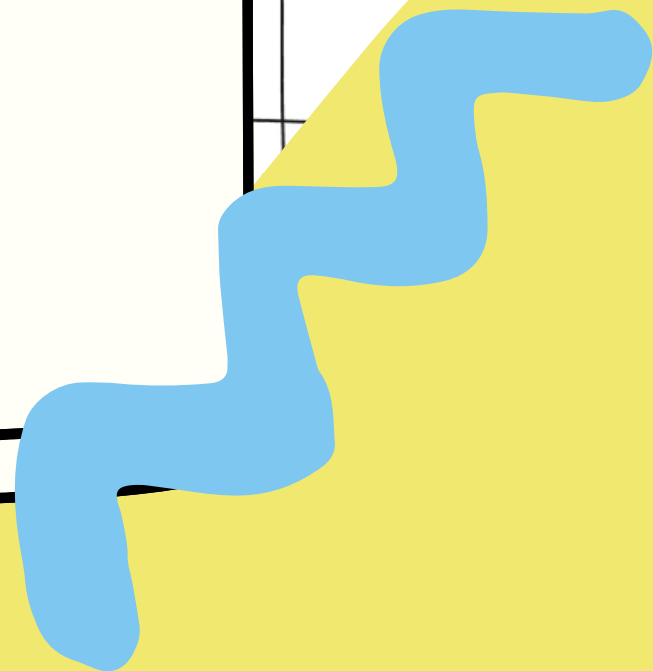
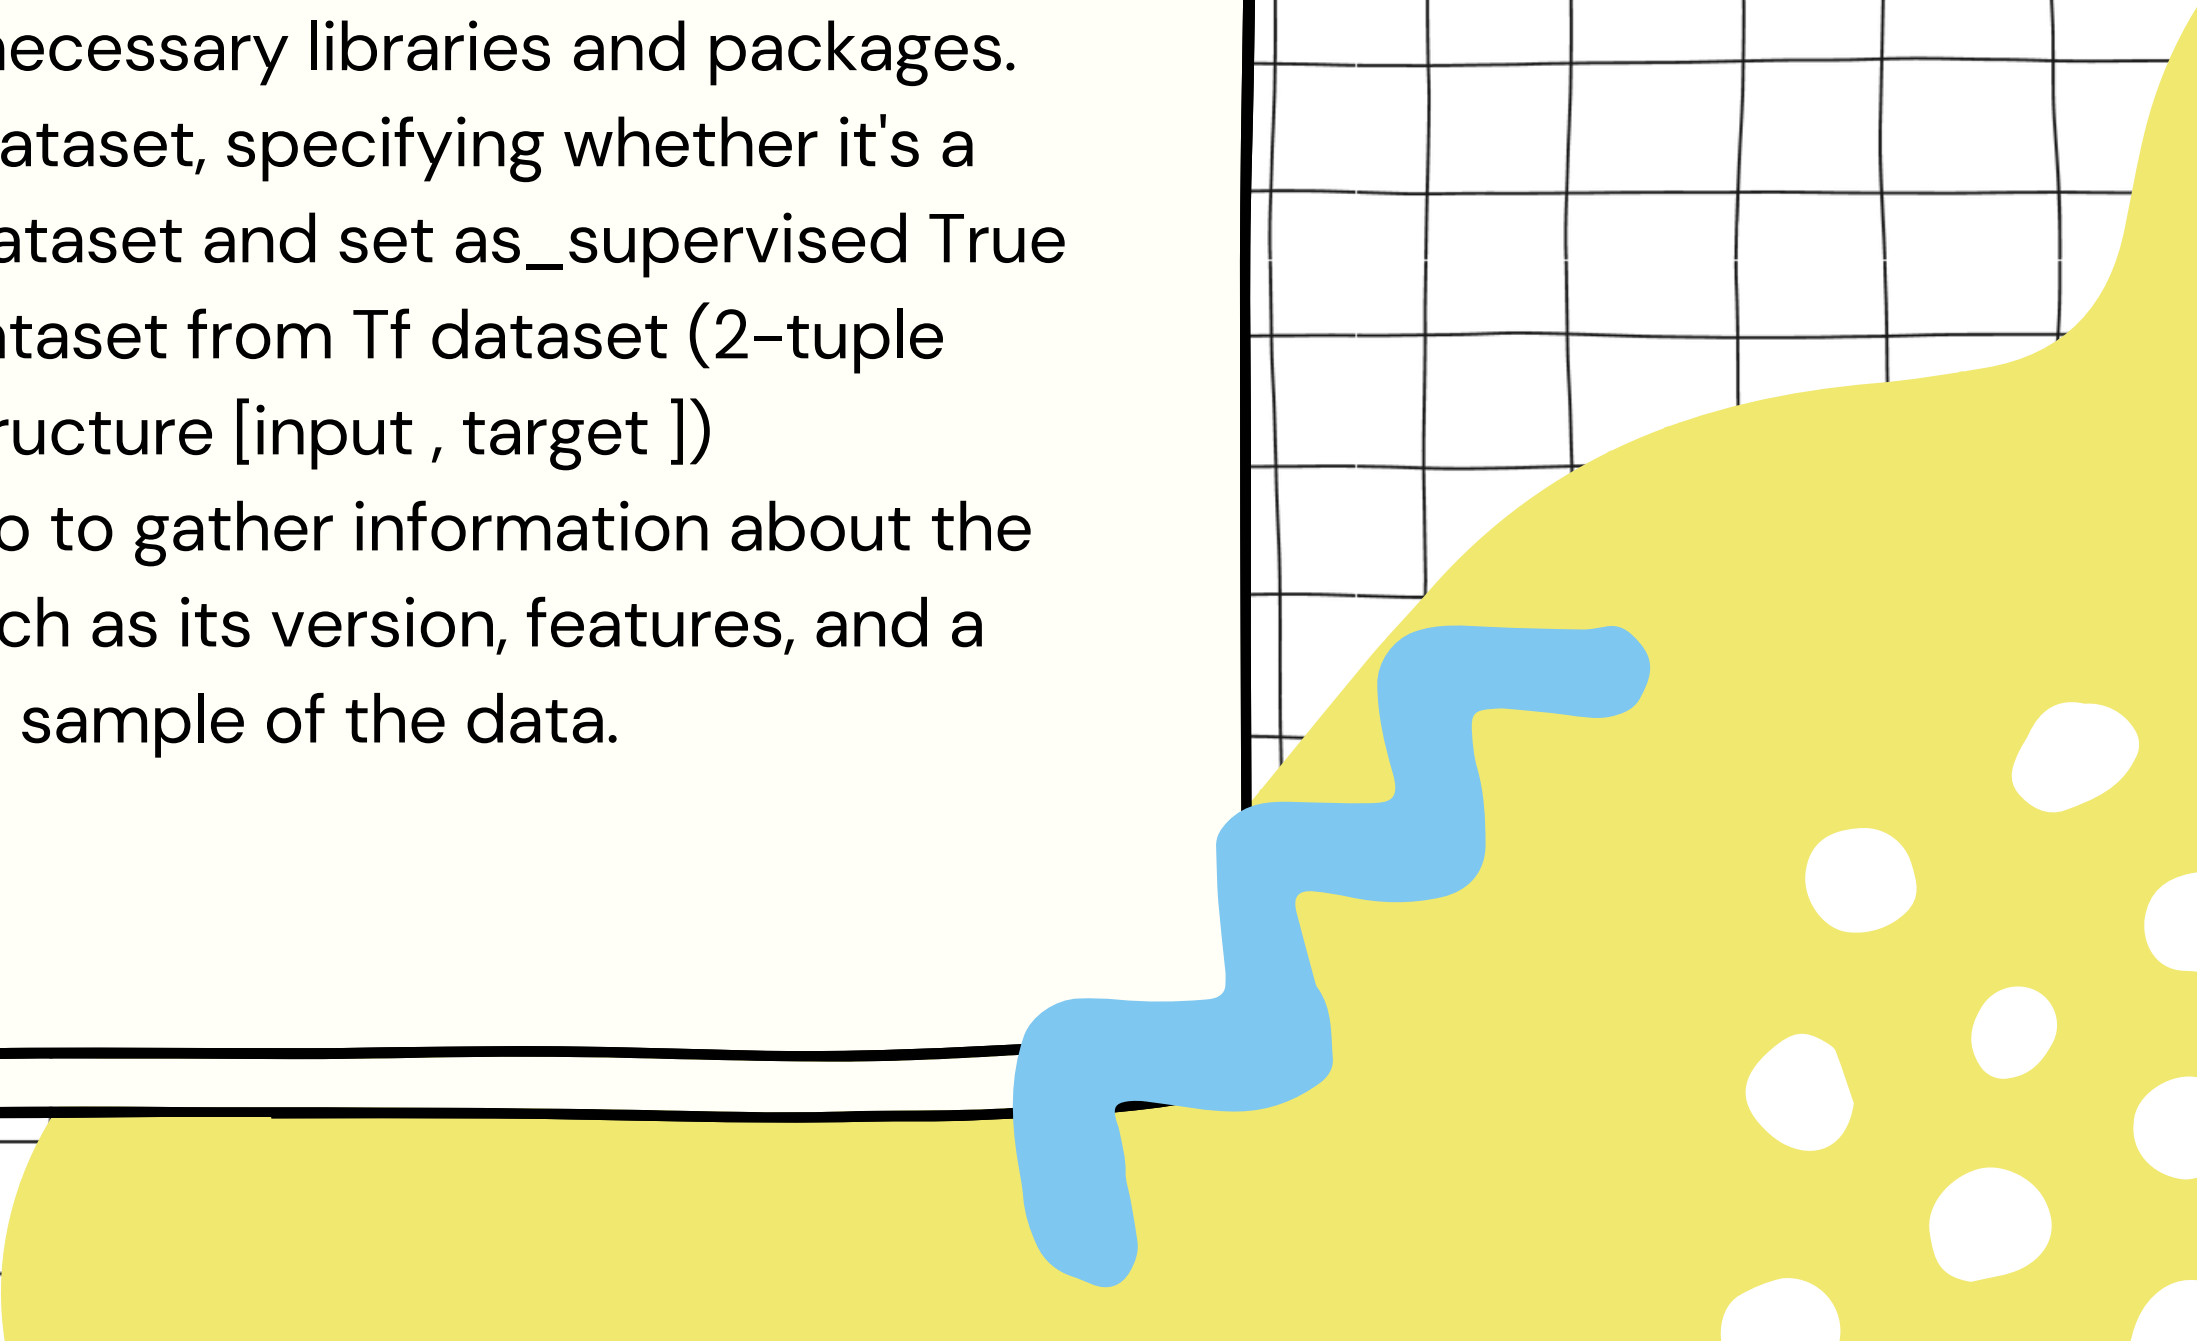
04

Model
Evaluation





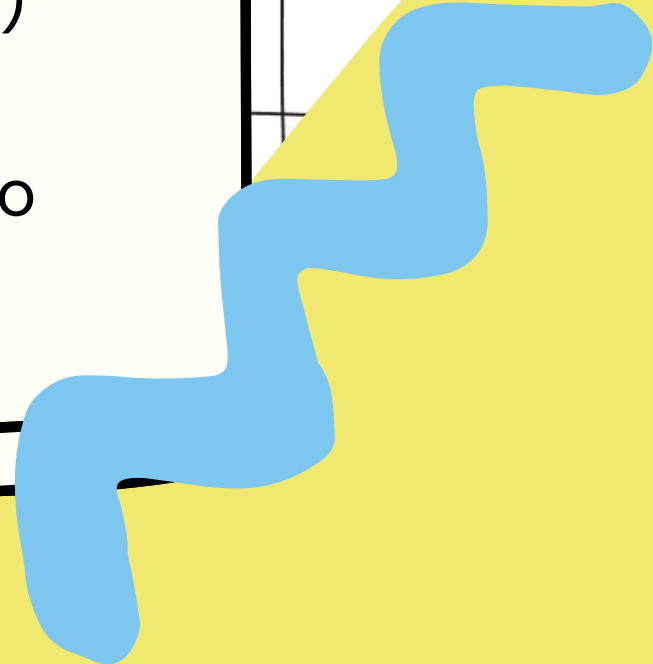
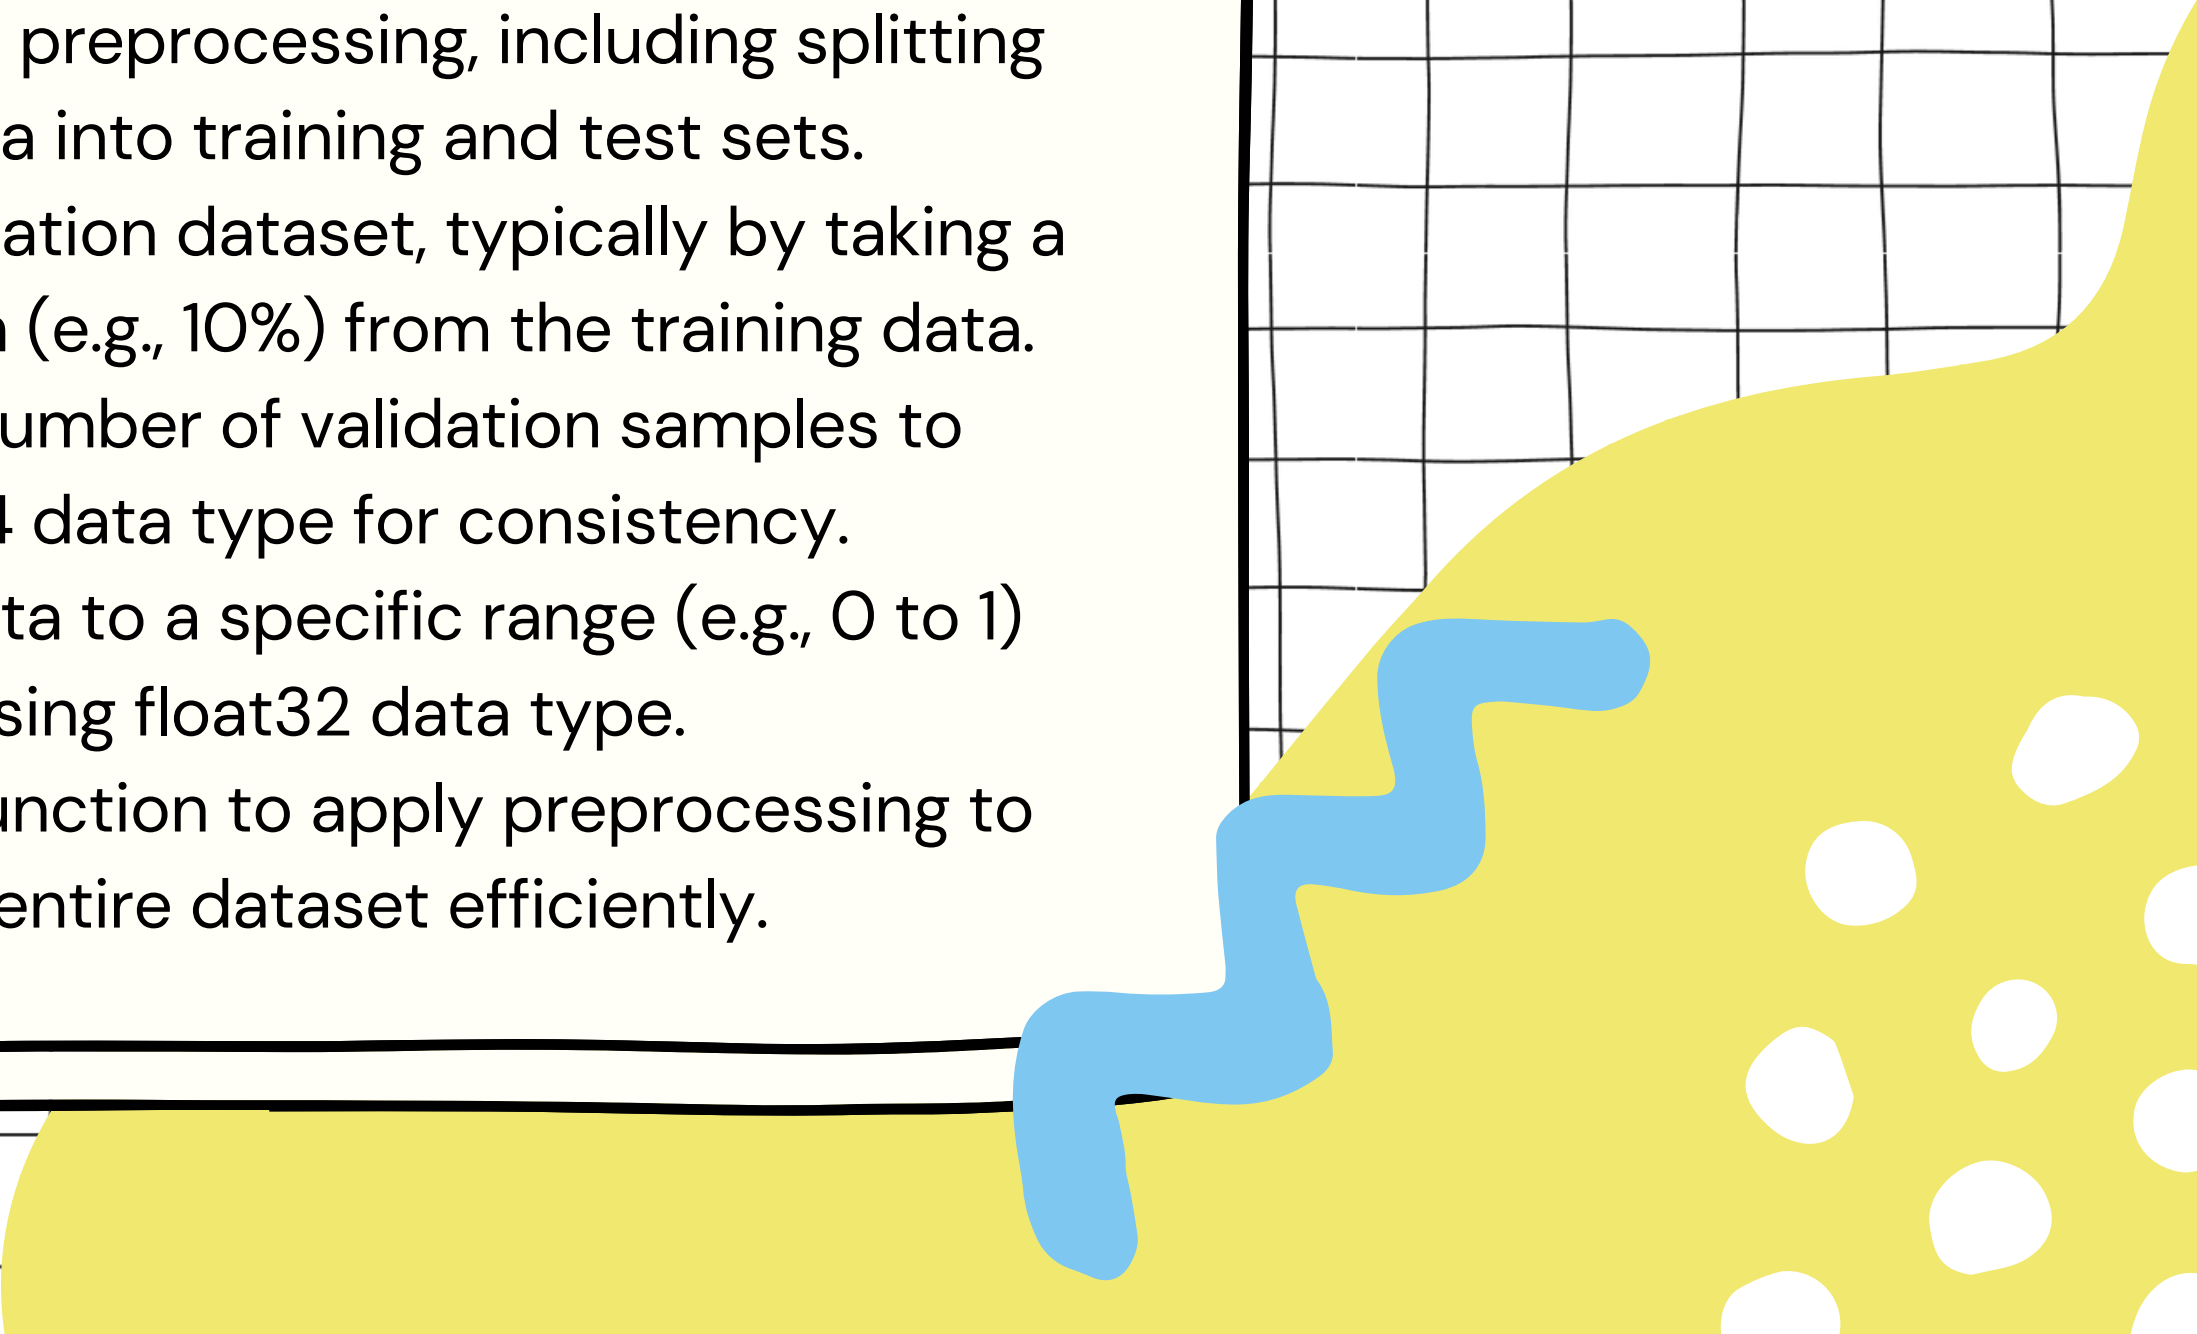
Data Import :

- Import the necessary libraries and packages.
 - Load the dataset, specifying whether it's a TensorFlow dataset and set `as_supervised True` : load a dataset from Tf dataset (2-tuple structure [input , target])
 - Use `with_info` to gather information about the dataset, such as its version, features, and a sample of the data.
- 
- 




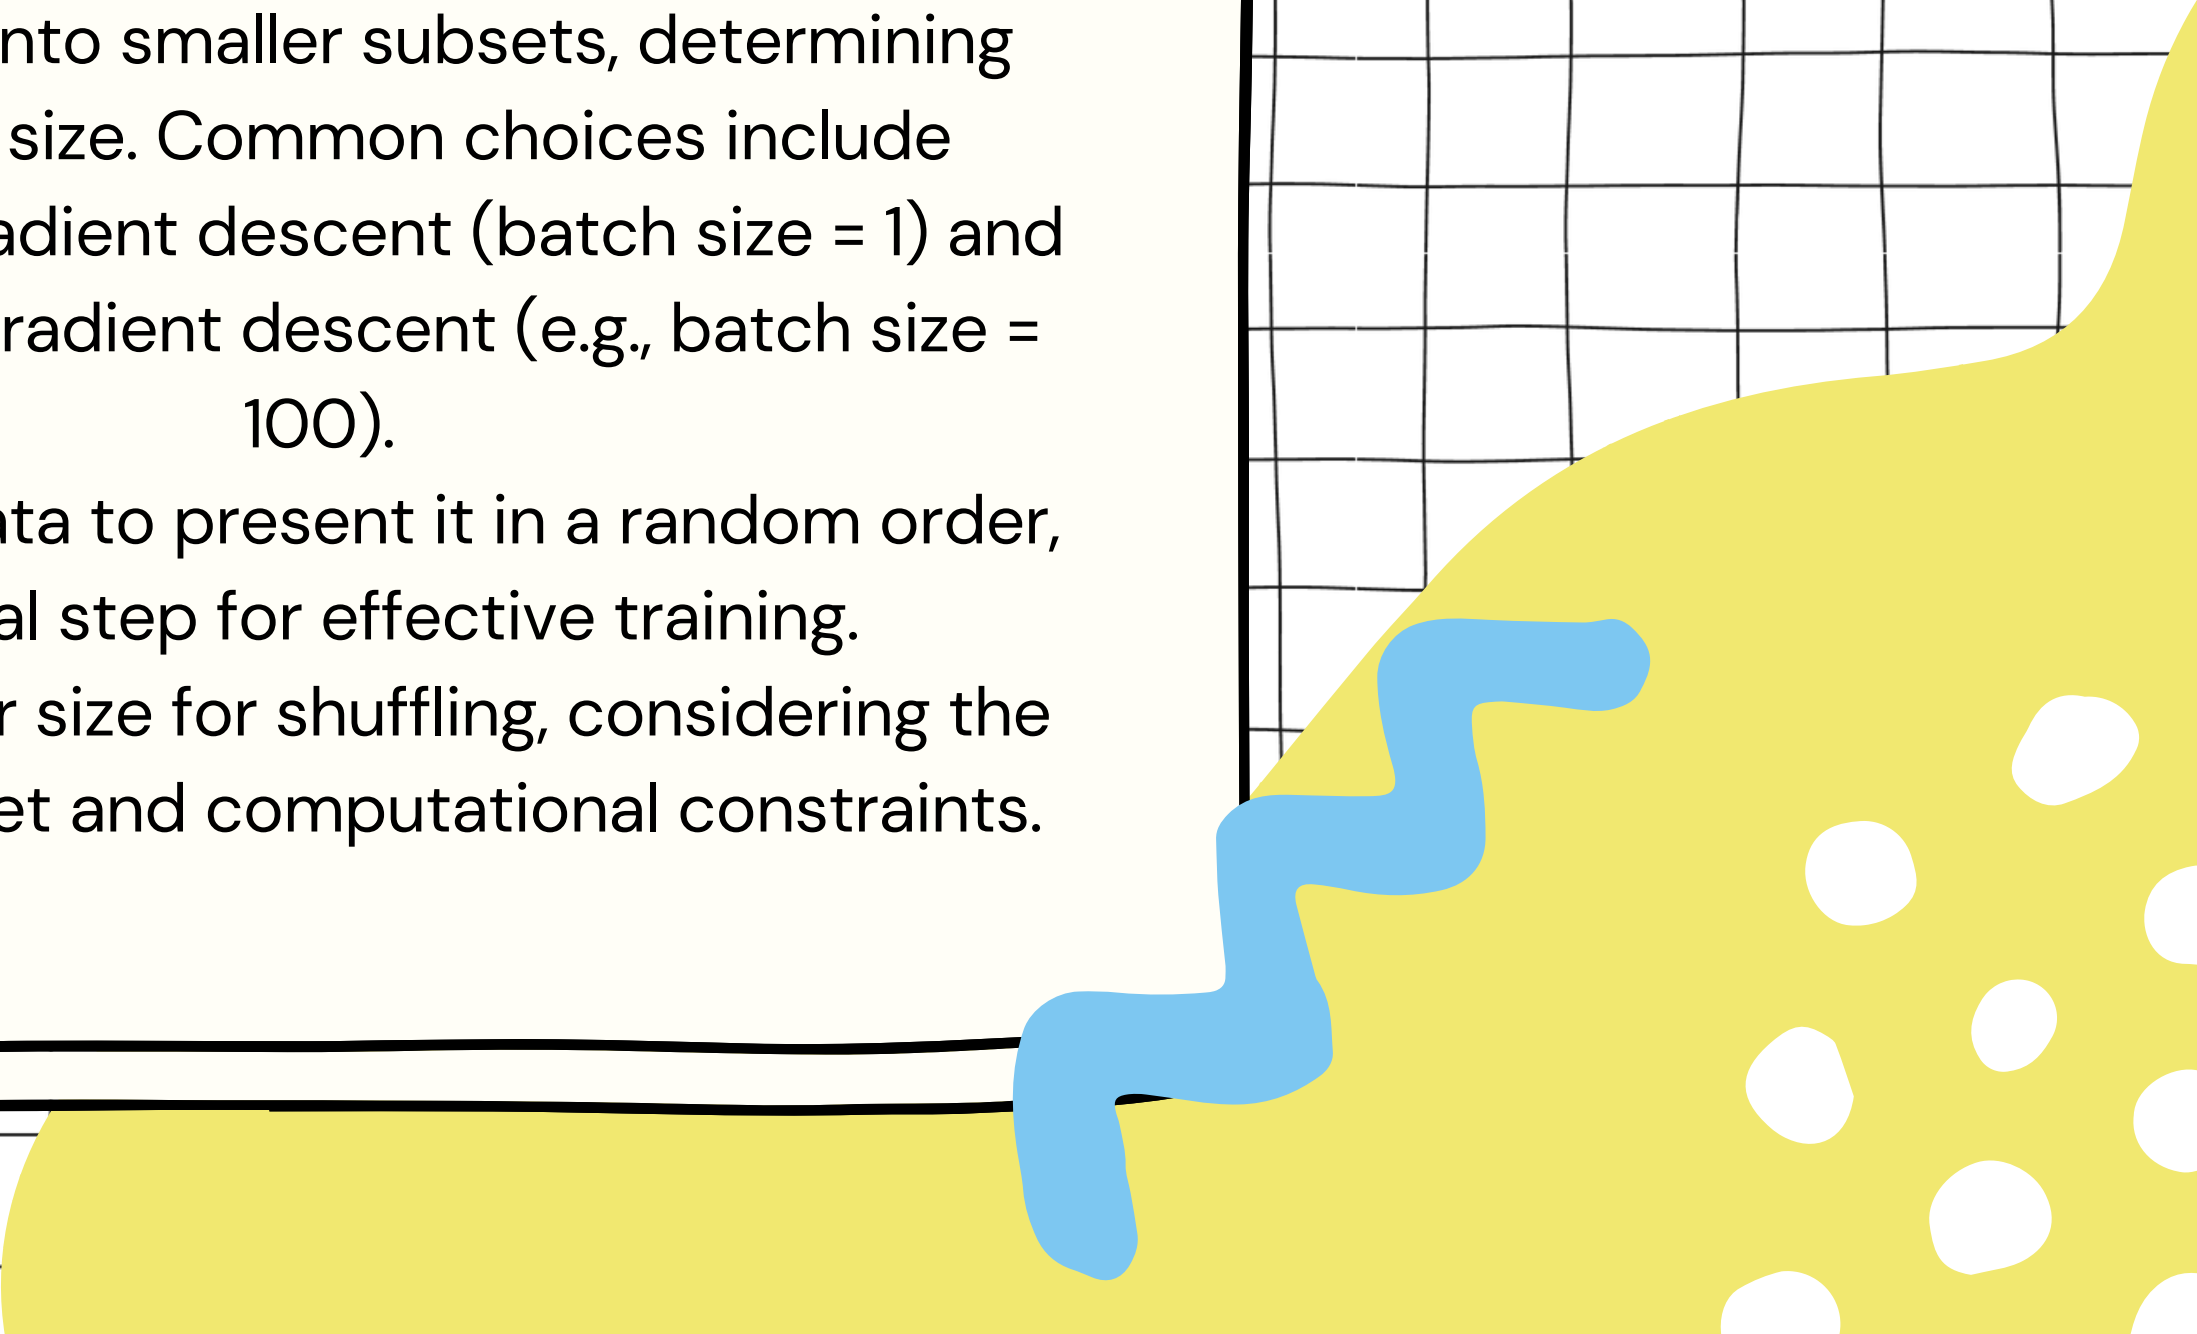
Data

Preprocessing:

- Perform data preprocessing, including splitting the data into training and test sets.
 - Create a validation dataset, typically by taking a small portion (e.g., 10%) from the training data.
 - Cast the number of validation samples to `tf.int64` data type for consistency.
 - Scale the data to a specific range (e.g., 0 to 1) using `float32` data type.
 - Use a map function to apply preprocessing to the entire dataset efficiently.
- 
- 


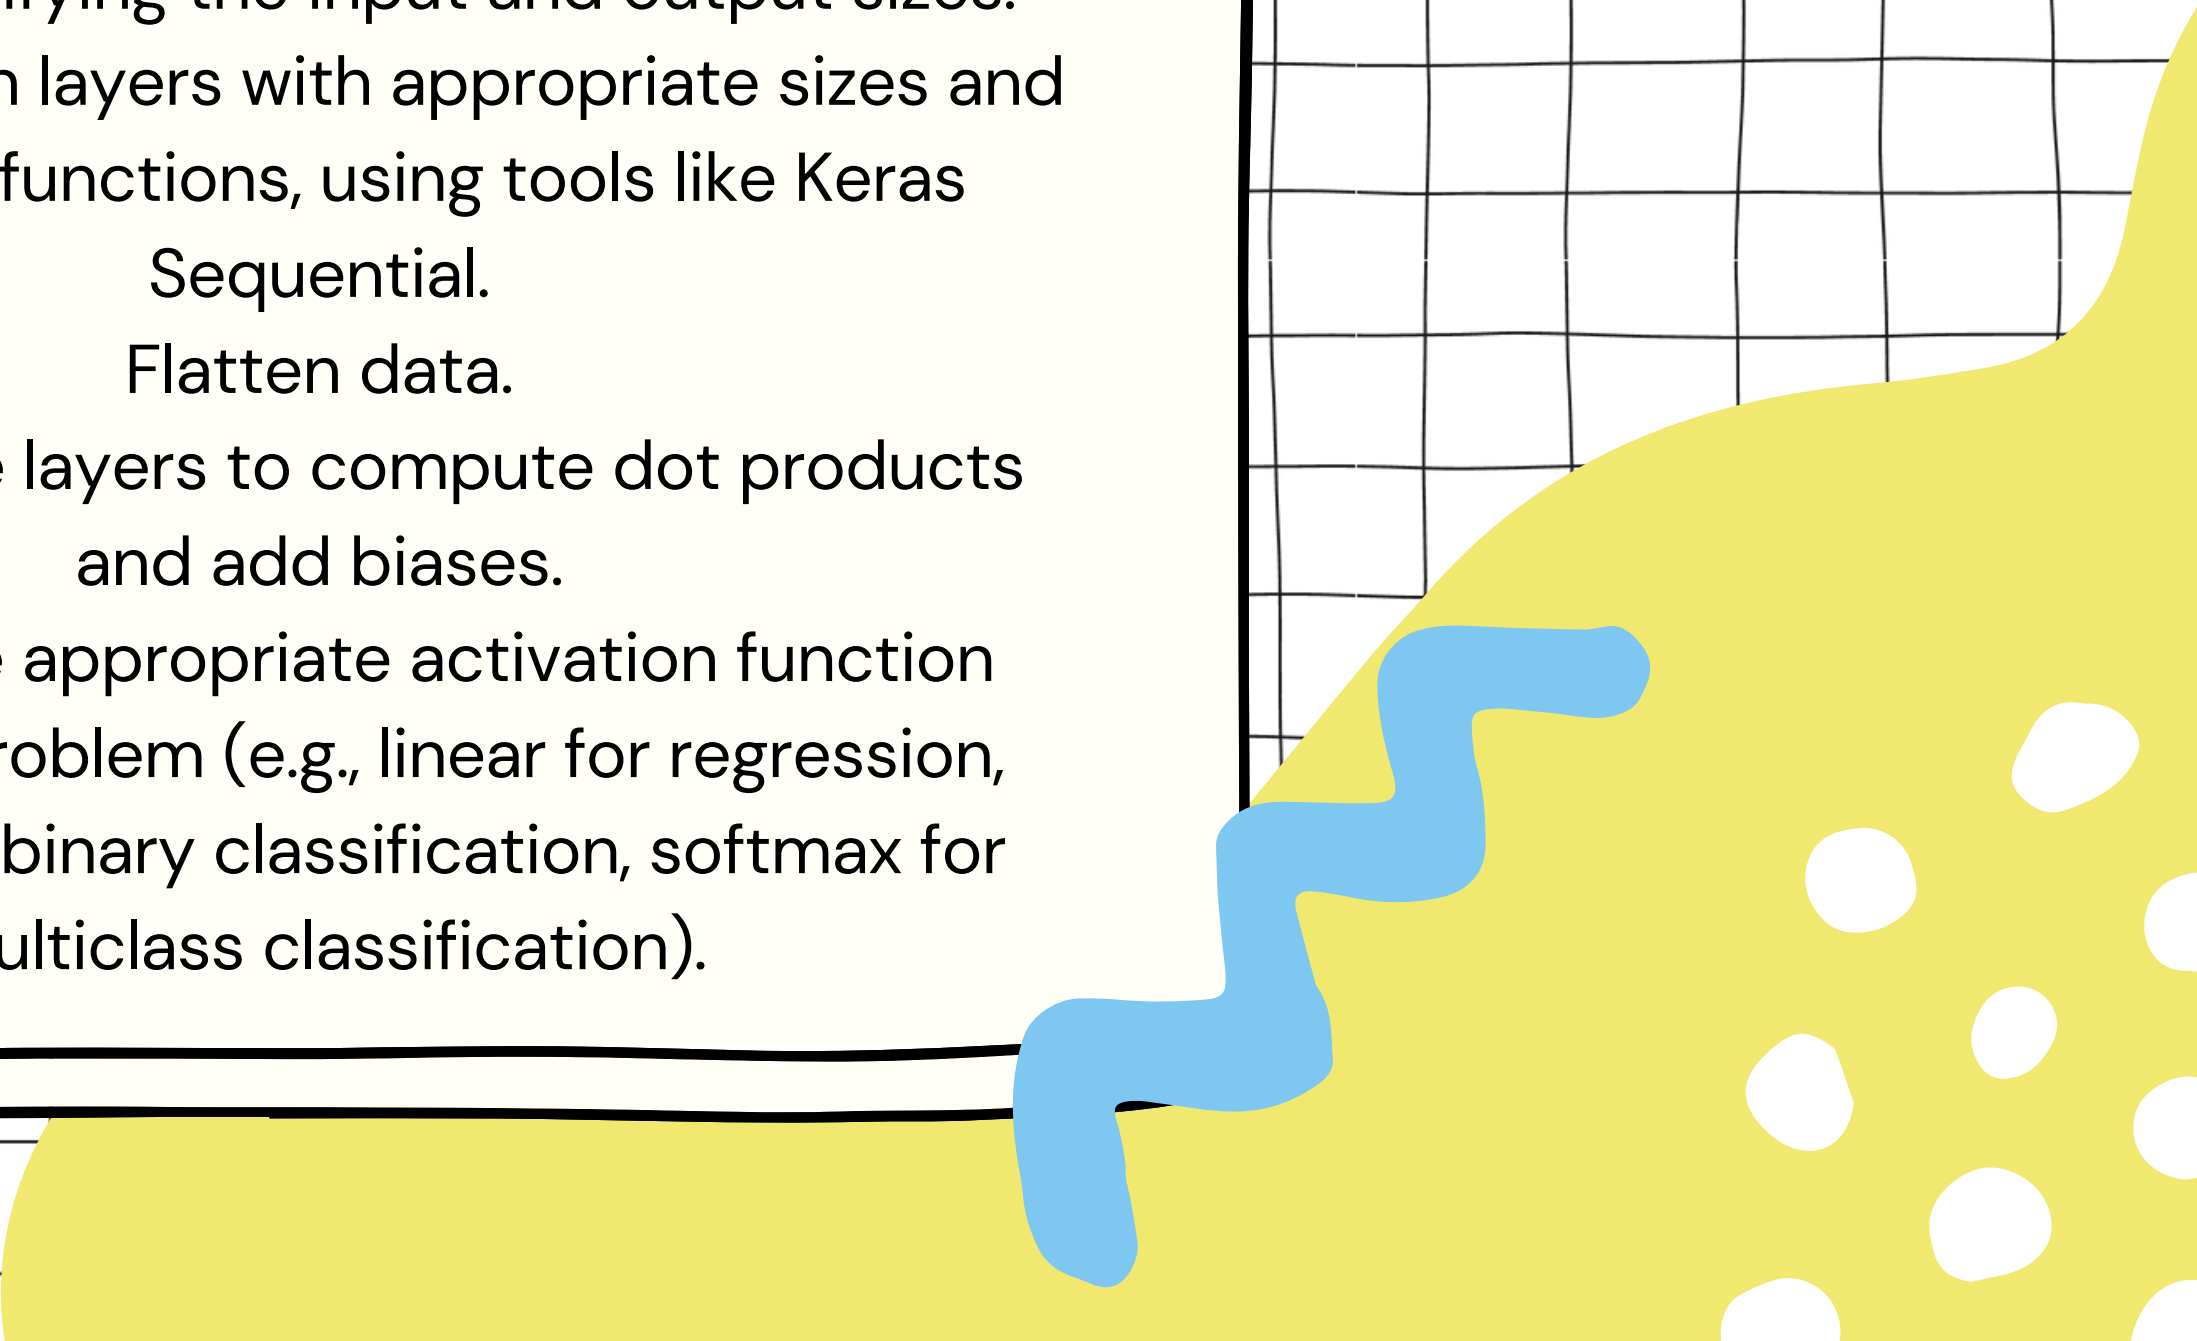


Batching and Shuffling:

- Batch data into smaller subsets, determining the batch size. Common choices include stochastic gradient descent (batch size = 1) and mini-batch gradient descent (e.g., batch size = 100).
 - Shuffle the data to present it in a random order, a crucial step for effective training.
 - Set the buffer size for shuffling, considering the size of dataset and computational constraints.
- 
- 



Model Outline:

- Define the architecture of machine learning model, specifying the input and output sizes.
 - Create hidden layers with appropriate sizes and activation functions, using tools like Keras Sequential.
 - Flatten data.
 - Apply dense layers to compute dot products and add biases.
 - Choose the appropriate activation function based on problem (e.g., linear for regression, sigmoid for binary classification, softmax for multiclass classification).
- 
- 



Optimize and choose the appropriate loss function :

- Configure the model for training by compiling it.
- Select an optimizer, such as Adam, for efficient gradient-based optimization.
- Choose a loss function suitable for your problem (e.g., binary cross-entropy, categorical cross-entropy).
- Define metrics to monitor model performance, often including accuracy.

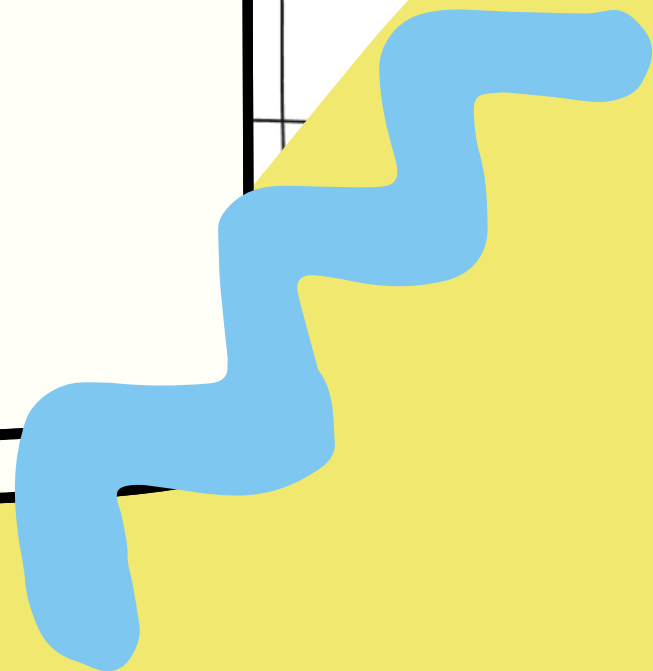
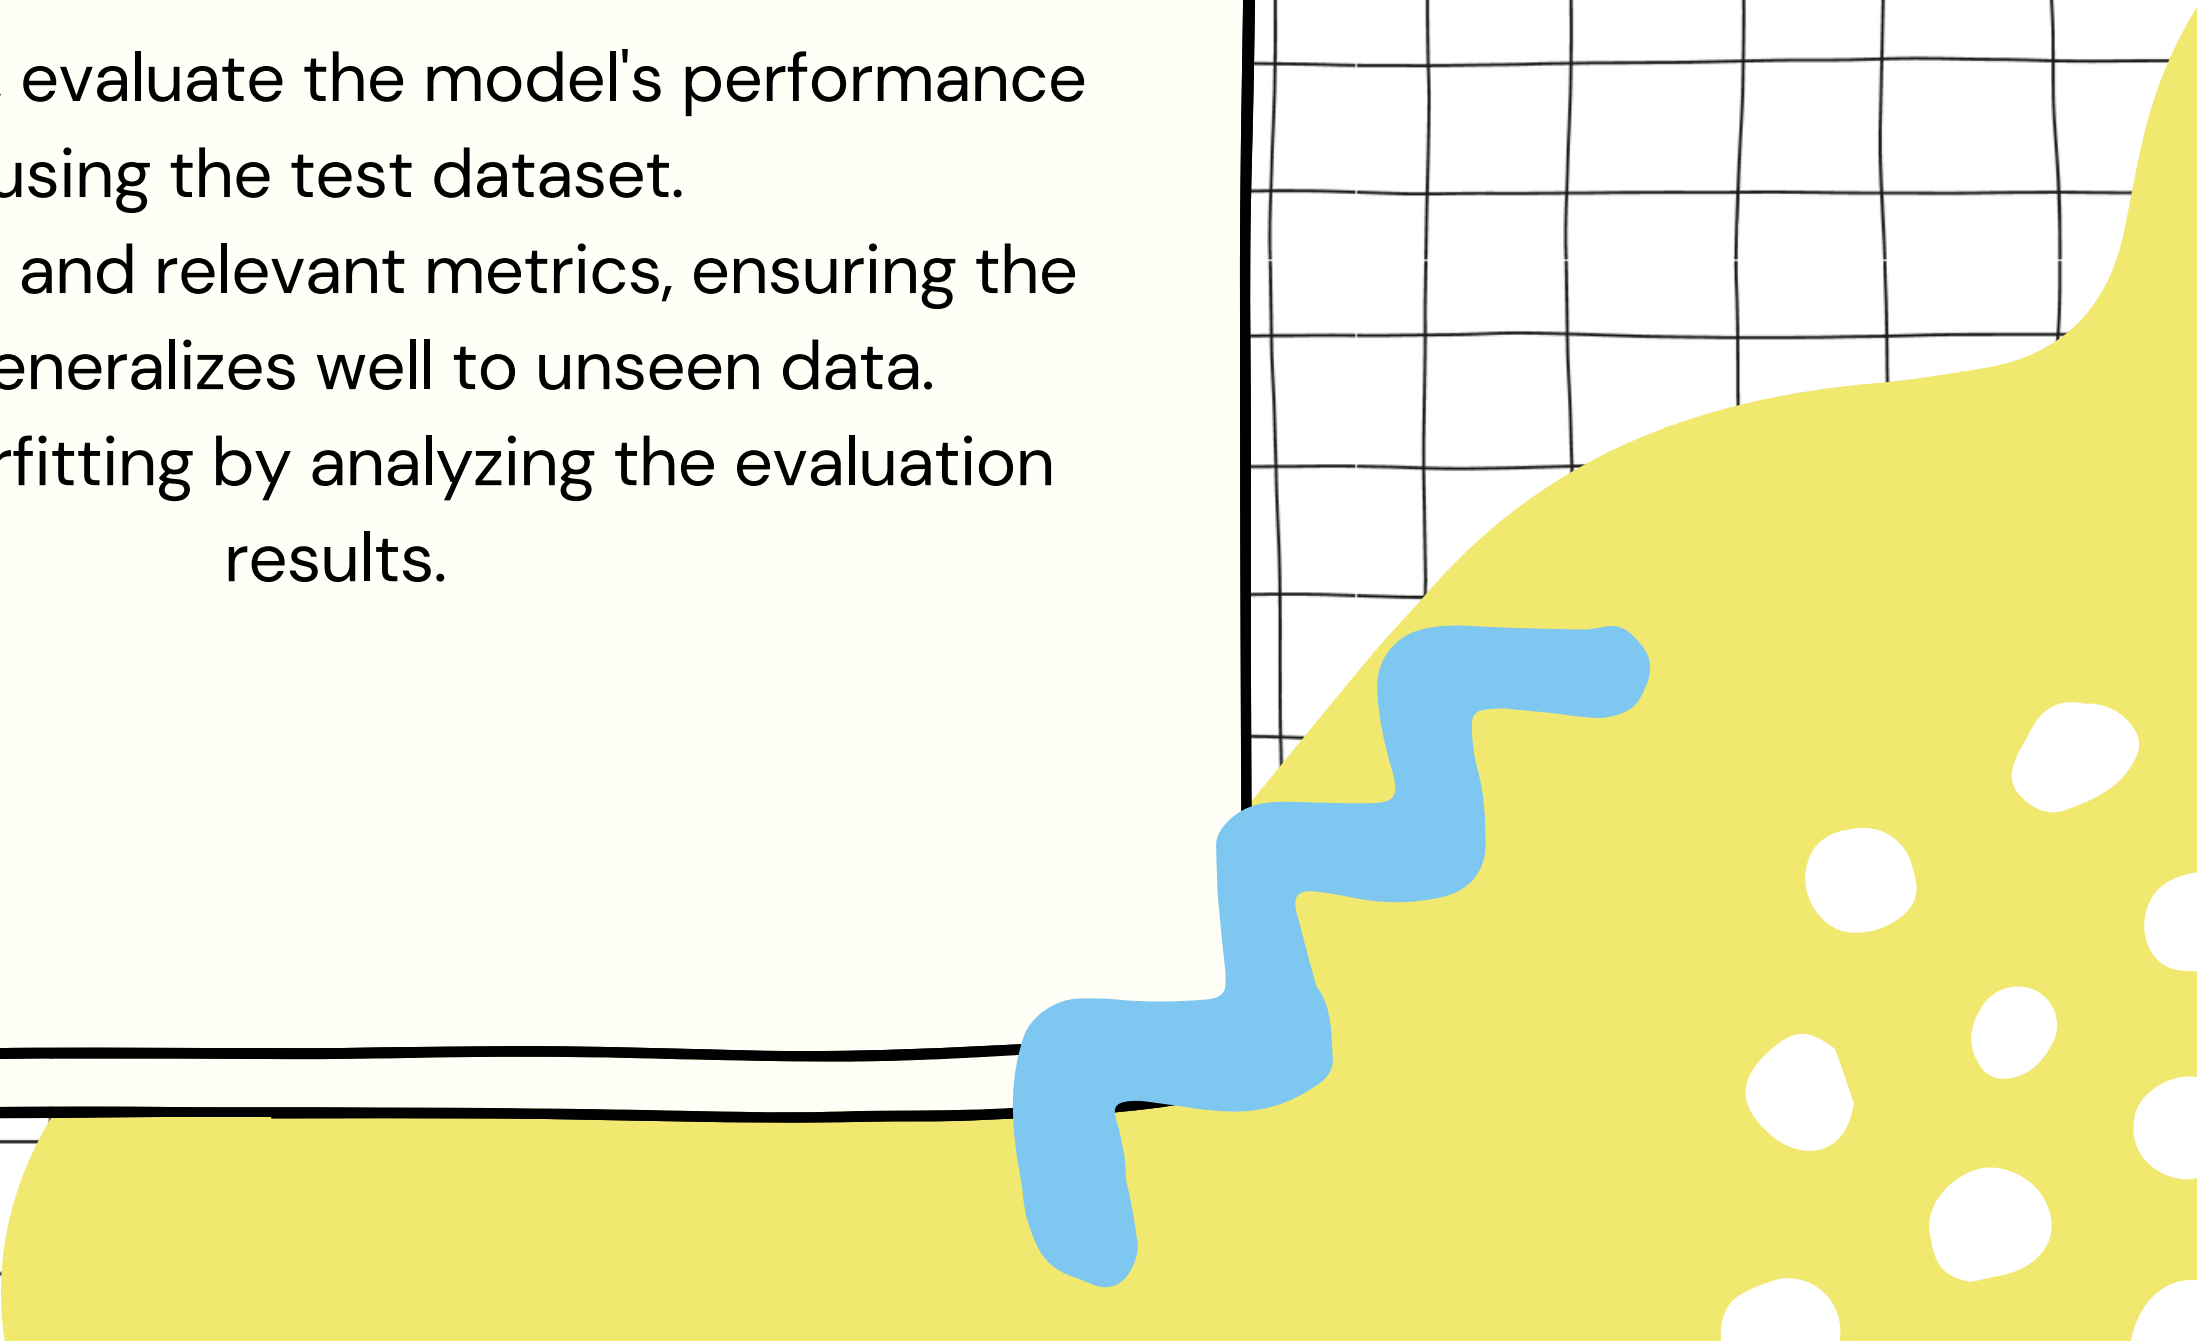


Training the Model & Improvements

- Specify the number of epochs for training, determining how many times the model will see the entire dataset.
- Train the model by fitting it to the training data, using the validation dataset to monitor performance.
- Adjust the model architecture as needed, including changing the number of hidden layers or their sizes, to improve performance.



Model Evaluation

- After training, evaluate the model's performance using the test dataset.
 - Measure loss and relevant metrics, ensuring the model generalizes well to unseen data.
 - Prevent overfitting by analyzing the evaluation results.
- 
- 



Thank you